

Prajwal Gowda

AI Engineer · Agentic AI · LLMOps · Multi-Agent Systems · RAG

LinkedIn GitHub prajwalgowdads2709@gmail.com +91 9448145109 Bengaluru, India

PROFILE

AI Engineer with 1.5+ years building agentic systems from 3-level multi-agent orchestration platforms to offline RAG pipelines achieving 97% answer accuracy. Delivered sub-5s inference latency on single-GPU hardware through full-stack LLM optimization, eliminating entire manual reporting workflows. Published ML researcher (MLIP-2025) with a 9.38 GPA and end-to-end ownership across the complete AI engineering stack.

TECHNICAL SKILLS

Agentic AI: LangGraph, LangChain, MCP, Multi-Agent Orchestration, Tool Use, Memory, Fallback Logic

LLM Infra: vLLM, AWQ, OpenAI API, LLM Inferencing, Langfuse Observability, LLMops, AIOps, LoRA / QLoRA Fine-tuning, LLM evals, Transformers

RAG & Retrieval: Vector DB, Semantic Chunking, Hybrid Retrieval, Reranking, OCR Metadata Extraction

Backend & Infra: Python, PostgreSQL, Redis, Docker, Linux, FastAPI, REST APIs, Git, Streamlit, Auto-Scaling, Rate Limiting

ML / DL: PyTorch, TensorFlow, Scikit-learn, CNN, RNN, FFNN, NLP, SMOTE, Transfer Learning, EDA, Feature Engineering

Cloud: GCP

EXPERIENCE

AI Engineer R&D

Aug 2025 - Present

Tejas Networks · Bengaluru · Full-time

- Architected a multi-agent AI from scratch, reducing **average p95 response latency by 70% to under 2 to 5 seconds on a single entry-level GPU** eliminating multi-GPU hardware cost by tuning vLLM inference, custom routing, Redis caching, auto-scaling, and Langfuse observability across a 3-level agent orchestration stack.
- Engineered **LangGraph state-machine-based dynamic report generation** and **on-demand dashboard/chart creation** from natural language queries, **eliminating manual reporting workflows with 100% workflow automation**.
- Delivered full project ownership in vanilla Python + selective LangChain/LangGraph, designed dynamic agent routing, persistent memory.
- Architected a **Config Assist for cross-vendor CLI translation** engine that translates network device configuring commands from other network vendors into Tejas Network's CLI, **eliminating manual lookup of configuration commands for network team**.

AI Engineer Intern

Jan 2025 - Aug 2025

Tejas Networks · Bengaluru

- Deployed custom offline RAG pipeline achieving **97% answer accuracy vs. .30% baseline vector search from +37 pp improvement** with sub-second retrieval serving SOPs, and tech docs with zero external API dependency by combining semantic chunking, OCR metadata extraction, embedding, and reranking.
- Improved retrieval precision significantly over baseline vector search by building custom retrievers with dynamic document chunking and multi-stage reranking, integrated as a dedicated retrieval agent in the multi-agent platform.

Data Analyst Intern

Oct 2024 - Nov 2024

Tata Steel · Jamshedpur · Remote · Vocational Training

- Conducted EDA on large-scale operational datasets to surface KPI anomalies and resource utilisation patterns, delivering findings as Tableau and Excel reports to the Data Science team.

PROJECTS

Personal Agentic AI Assistant

2026 - Present

Stack: Python · AirLLM · LangGraph · MCP · Milvus · TTS/STT · Vision LLM · Various REST APIs · Graph Memory

- Built 2-layer orchestrated assistant: local Qwen3.6 via AirLLM for sensitive queries + cloud LLMs, with MCP, search agents, web scraping with docking parser, Neo4j graph-based memory with obsidian support, hybrid RAG, and WhatsApp interface.
- Integrated TTS/STT, Vision LLM, MS365 (PPT/DOCX/Excel), reminders, and interactive dashboards into a unified personal AI stack.

Deep-Learning Based Early Diagnosis of ASD [Published · MLIP-2025]

2024 - 2025

Stack: Python · TensorFlow · PyTorch · Scikit-learn · Nilearn · PyCaret · OpenCV · SMOTE

- Built Multi-modal DL pipeline: hyperparameter-tuned SVC for gene expression (**87% accuracy, AUC-ROC 0.927**), Custom-FFNNs for fMRI connectivity, Custom-CNNs for facial imagery.
- Applied SMOTE + augmentation + cross-modal tuning to address class imbalance — enabling reliable early-stage ASD detection.

AquaVision: Real-Time Marine Species Detection & Underwater Waste Detection

2024

Stack: Python · TensorFlow · EfficientNetB7 · OpenCV · YOLOv8 · TensorFlow · Pandas · NumPy

- Real-time sea creature detection & classification using EfficientNetB7 with full ETL/EDA pipeline for marine biodiversity research.
- YOLOv8-based real-time oceanic trash detection & classification system for environmental conservation.

EDUCATION

B.E. Artificial Intelligence & Machine Learning

2021 - 2025

Bangalore Institute of Technology (VTU) · Bengaluru · GPA: 9.38 / 10

Paper Published: Deep-Learning Based Early Diagnosis of Autism Spectrum Disorder Using Facial Features, Neuroimaging, and Blood Biomarkers - Hinweis International Conference - Machine Learning & Intelligent Processing (MLIP-2025)

Extracurricular: President of Rotaract Club of BIT - Drove **70% membership growth** and organised 37+ events including workshops, social impact programs, and skill development initiatives.